

Shared-memory programming : OpenMP (I)

ING2-GSI-MI Architecture et Programmation Parallèle

Juan Angel Lorenzo del Castillo
juan-angel.lorenzo-del-castillo@cyu.fr

CY Cergy Paris Université
2023-2024



Table of Contents

1 Introduction

2 OpenMP

Table of Contents

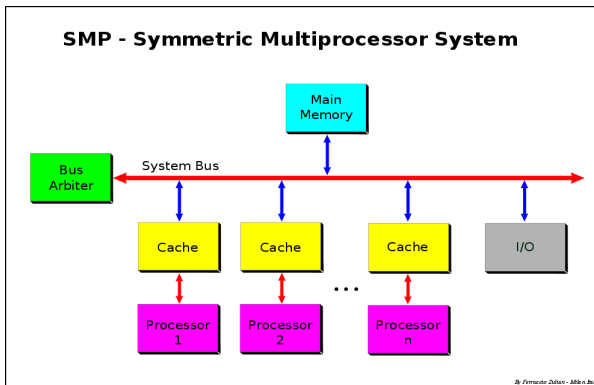
1 Introduction

2 OpenMP

Shared-memory programming

Parallel programming

A program is split in processes or threads that cooperate and coordinate together.



Source : Wikipedia

Shared-memory programming

Parallel programming

A program is split in processes or threads that cooperate and coordinate together.

Shared-memory programming

- Based on the notion of *threads*
- Coordination and cooperation are done by writing and reading shared variables and by using synchronisation variables
 - ▶ **Low level:** barriers, locks, critical sections, semaphores...
 - ▶ **High level:** compilation **directives** (*OpenMP*).

Shared-memory programming

Parallelisation with directives

- **Directive:** A special line of source code with meaning only to certain compilers. It is distinguished by a sentinel at the start of the line.
 - ▶ The developer adds the compilation directives to the code.
 - ▶ The compiler will convert them into library calls.
- The original code is not modified.
- Portable, fast (but perhaps not too flexible nor efficient).
- Example: `#pragma omp`

Shared-memory programming


OpenMP is based on the notion of ***threads***

Reminder: Threads.

- Threads are like processes, except that threads can share memory with each other (as well as having private memory)
- Shared data can be accessed by all threads
- Private data can only be accessed by the owning thread
- Different threads can follow different flows of control through the same program
- Usually run one thread per processor (but could be more)
- Definitions :
 - ▶ *Thread team*: set of threads which cooperate on a task.
 - ▶ *Master thread*: is responsible for coordinating the team.

Shared-memory programming

Thread communication:



Source : EPCC

Table of Contents

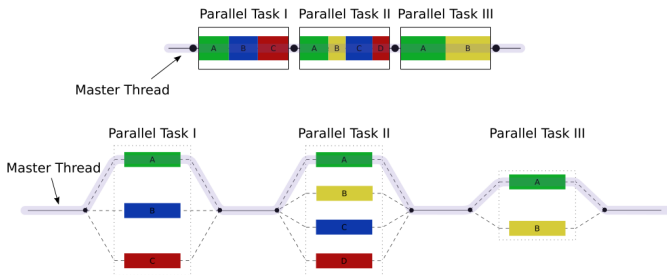
1 Introduction

2 OpenMP

Presentation

OpenMP (*Open MultiProcessing*)

- Set of directives, libraries and variables available for Fortran, C and C++.
- It is the standard on shared-memory programming.
- References:
 - ▶ www.openmp.org
 - ▶ www.compunity.org
- Based on the `fork-join` execution model.



Source : Wikipedia

OpenMP directives

Most relevant OpenMP directives

- Parallel regions construction
 - ▶ `parallel`
- Work sharing
 - ▶ `for`, `sections`, `single`.
- Synchronisation
 - ▶ `master`, `critical`, `atomic`, `barrier`.
- Task management
 - ▶ `task`, `taskwait`.

There are more...

OpenMP directives

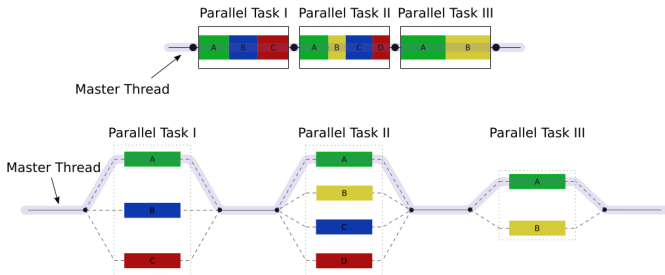
Most relevant OpenMP directives

- Parallel regions construction
 - ▶ `parallel`
- Work sharing
 - ▶ `for`, `sections`, `single`.
- Synchronisation
 - ▶ `master`, `critical`, `atomic`, `barrier`.
- Task management
 - ▶ `task`, `taskwait`.

There are more...

Today's class

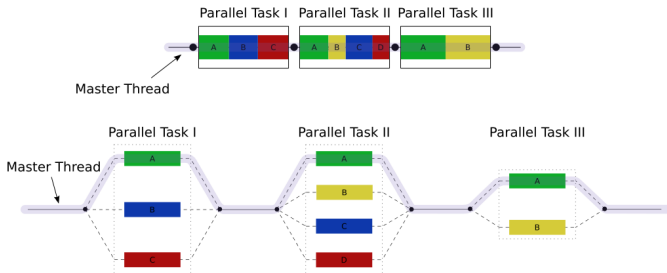
Parallel regions



Source : Wikipedia

- A program begins execution on a single thread (the *master* thread).
- When the first parallel region is encountered, the master thread creates a team of threads (fork/join model).
- Every thread executes the statements which are inside the parallel region.
- At the end of the parallel region, the master thread waits for the other threads to finish, and continues executing the next statements.

Parallel regions



Source : Wikipedia

- Inside a parallel region, variables can be either shared (all threads see the same copy) or private (each thread has its own copy).
- All threads see the same copy of shared variables.
- All threads can read or write shared variables.
- Each thread can have its own private variables, invisible to the other threads.
- A private variable can only be read or written by its owner thread.
- **Implicit barrier** at the end of the parallel region.

Parallel regions

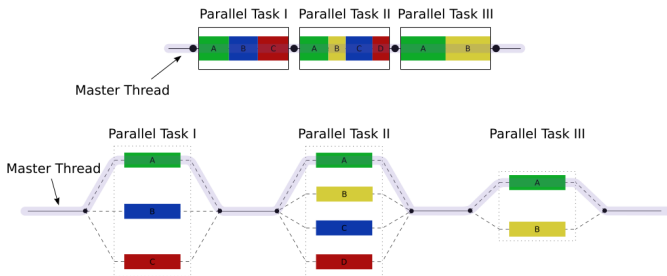
OpenMP example:

```
1 # include <omp.h>
2 # include <stdio.h>
3
4 int main () {
5
6     int sharedvar = 1;
7     printf(" Hola! I am the master thread :-) \n");
8
9     # pragma omp parallel
10    {
11        int id = omp_get_thread_num();
12        printf("I am thread %d. sharedvar = %d\n",id , sharedvar);
13    }
14
15    printf(" The master thread ends... \n");
16    return 0;
17 }
```

Compilation:

```
gcc -fopenmp -Wall -Wextra -o example01 example01.c
```

Parallel regions



Source : Wikipedia

- Note that the variables defined in the sequential section are **shared** inside the parallel region.
- Variables defined inside the parallel region are **private** to each thread.
- If the parallel region contains function calls, the local variables in the function will be **private**.

Setting/getting the number of threads

Setting the number of threads:

- As an environment variable (`bash/ksh`):
 - ▶ `export OMP_NUM_THREADS=8`

Setting/getting the number of threads

Setting the number of threads:

- As an environment variable (bash/ksh):

▶ `export OMP_NUM_THREADS=8`

1 `#!/bin/bash`

2

3 `export OMP_NUM_THREADS=4`

4 `./example01`

Setting/getting the number of threads

Setting the number of threads:

- As an environment variable (bash/ksh):
 - ▶ `export OMP_NUM_THREADS=8`
- In the program code:
 - ▶ `void omp_set_num_threads(num_threads);`

Setting/getting the number of threads

Setting the number of threads:

- As an environment variable (bash/ksh):

- ▶ `export OMP_NUM_THREADS=8`

- In the program code:

- ▶ `void omp_set_num_threads(num_threads);`

```
1 # include <omp.h>
2 # include <stdio.h>
3
4 int main () {
5     omp_set_num_threads(8);
6
7     printf(" Hola! I am the master thread :-) \n");
8
9     # pragma omp parallel
10    {
11        int id = omp_get_thread_num();
12        printf("I am thread %d \n",id);
13    }
14
15    printf(" The master thread ends... \n");
16    return 0;
17 }
```

Setting/getting the number of threads

Setting the number of threads:

- As an environment variable (bash/ksh):
 - ▶ `export OMP_NUM_THREADS=8`
- In the program code:
 - ▶ `void omp_set_num_threads(num_threads);`
- Inside a parallel region:
 - ▶ `#pragma omp parallel num_threads(num_threads)`

Setting/getting the number of threads

Setting the number of threads:

- As an environment variable (bash/ksh):
 - ▶ `export OMP_NUM_THREADS=8`
 - In the program code:
 - ▶ `void omp_set_num_threads(num_threads);`
 - Inside a parallel region:
 - ▶ `#pragma omp parallel num_threads(num_threads)`
- ```
1 # include <omp.h>
2 # include <stdio.h>
3
4 int main () {
5 printf(" Hola! I am the master thread :-) \n");
6
7 # pragma omp parallel num_threads(8)
8 {
9 int id = omp_get_thread_num();
10 printf("I am thread %d \n",id);
11 }
12
13 printf(" The master thread ends... \n");
14 return 0;
15 }
```

# Setting/getting the number of threads

## Getting the number of threads (inside the program):

- Number of **running** threads in the team.

- ▶ `int omp_get_num_threads();`

- Thread id

- ▶ `int omp_get_thread_num();`

- ▶ Between 0 and `omp_get_num_threads() - 1`

# Setting/getting the number of threads

## Getting the number of threads (inside the program):

- Number of **running** threads in the team.

- ▶ `int omp_get_num_threads();`

- Thread id

- ▶ `int omp_get_thread_num();`

- ▶ Between 0 and `omp_get_num_threads() - 1`



# Setting/getting the number of threads

## Getting the number of threads (inside the program):

- Number of **running** threads in the team.

- ▶ `int omp_get_num_threads();`

- Thread id

- ▶ `int omp_get_thread_num();`

- ▶ Between 0 and `omp_get_num_threads() - 1`

```
1 # include <omp.h>
2 # include <stdio.h>
3
4 int main () {
5
6 omp_set_num_threads(8);
7 printf("How many threads (master): %d \n", omp_get_num_threads());
8
9 # pragma omp parallel num_threads(4)
10 {
11 int id = omp_get_thread_num();
12 printf("I am thread %d \n", id) ;
13 printf("How many threads in the parallel region (%d): %d \n\n", id ,
14 omp_get_num_threads());
15 }
16 printf(" The master thread ends... \n");
17 return 0;
18 }
```

# Setting/getting the number of threads

## Getting the number of threads (inside the program):

- Number of **running** threads in the team.
  - ▶ `int omp_get_num_threads();`
- Thread id
  - ▶ `int omp_get_thread_num();`
  - ▶ **Between 0 and** `omp_get_num_threads() - 1`

## Others :

- `int omp_get_max_threads();`
- `int omp_get_num_procs();`
- `int omp_in_parallel();`
- There are more...

# Parallel region

## Syntax:

```
#omp parallel [clause [clause]]
{
 // Parallel region
}
```

## Clauses allowed in a parallel region

- `if(logical expression)`
- `num_threads(integer)`
- `private(list of variables)`
- `shared(list of variables)`
- `default(shared | none)`
- `firstprivate(list of variables)`
- `reduction(operator: list of variables)`
- `copyin(list of variables)`

# Parallel region

## Syntax:

```
#omp parallel [clause [clause]]
{
 // Parallel region
}
```

## Clauses allowed in a parallel region

- `if(logical expression)`
- `num_threads(integer)`
- `private(list of variables)`
- `shared(list of variables)`
- `default(shared | none)`
- `firstprivate(list of variables)`
- `reduction(operator: list of variables)`
- `copyin(list of variables)`

# Parallel region

## #omp parallel if(logical expression)

- Creation of a thread team when the logical expression is different from 0. Otherwise, the execution will be sequential.

```
1 # include <omp.h>
2 # include <stdio.h>
3 # include <stdlib.h>
4
5 int main (int argc, char *argv[]) {
6
7 omp_set_num_threads(4);
8
9 if (argc != 3){
10 printf("Wrong number of parameters\n");
11 exit(0);
12 }
13
14 int a = atoi(argv[1]);
15 int b = atoi(argv[2]);
16
17 # pragma omp parallel if(a > b)
18 {
19 int id = omp_get_thread_num();
20 printf("I am thread %d\n", id);
21 }
22
23 return 0;
24 }
```

# Parallel region

## **#omp parallel num\_threads(integer)**

- Creation of a thread team that will run the parallel region.
- Already seen.

# Parallel region

## #omp parallel private(list of variables)

- The variables in the list are declared private for every thread.
  - ▶ Every thread will have a local copy of the variable.
  - ▶ The variables **are undefined** before and after the parallel region.

## #omp parallel shared(list of variables)

- The variables in the list are shared by all threads.
  - ▶ There is a danger of concurrent access.
  - ▶ Shared variables are used when:
    - There are read-only variables that we want to share.
    - Each thread accesses a different part of the variable (e.g. arrays).
    - We want to communicate a value to all threads.

## #omp parallel default(shared | none)

- All variables in the parallel region will be shared (`shared`) or undefined (`none`).

# Parallel region

## Examples (I) :

```
1 # include <omp.h>
2 # include <stdio.h>
3
4
5 int main () { // This code is wrong!!
6 omp_set_num_threads(4);
7
8 int a = 2;
9 int b = 3;
10
11 #pragma omp parallel private(a) shared(b)
12 {
13 int id = omp_get_thread_num();
14 #pragma omp critical //Ignore this pragma for now
15 {
16 printf("Thread %d. Before modification a: %d, b: %d\n",id ,a,b);
17
18 a++;
19 b++;
20
21 printf("Thread %d. After modification a: %d, b: %d\n",id ,a,b);
22 }
23 }
24
25 printf("Master thread. a: %d. b: %d\n",a,b);
26 return 0;
27 }
```



# Parallel region

## Examples (II) :

```
1 # include <omp.h>
2 # include <stdio.h>
3
4 int main () {
5 omp_set_num_threads(4);
6
7 int nthreads, tid;
8
9 #pragma omp parallel private(nthreads, tid)
10 {
11 tid = omp_get_thread_num();
12 printf("Hola! I am thread = %d\n", tid);
13
14 if (tid == 0) // If I am the master
15 {
16 nthreads = omp_get_num_threads();
17 printf("Number of threads = %d\n", nthreads);
18 }
19 }
20 }
```

# Parallel region

## Examples (III) :

```

1 #include <omp.h>
2 #include <stdio.h>
3 #define N 12
4
5 int fillArray(int, int *);
6
7 int main () {
8 omp_set_num_threads(4);
9 int id, i;
10 int A[N];
11 for (i=0; i<N; i++) A[i] = 12;
12
13 #pragma omp parallel private(id) shared(A)
14 {
15 id=omp_get_thread_num();
16 fillArray(id, A);
17 }
18
19 printf(" | ");
20 for (i=0; i<N; i++)
21 {
22 printf(" %d | ", A[i]);
23 }
24 printf("\n");
25
26 return 0;
27 }
```

```

27
28 int fillArray(int tid, int *array)
29 {
30 int begin, end, nth, chunk;
31 nth = omp_get_max_threads();
32 chunk = N/nth;
33 //integer division (on purpose)
34
35 printf("Thread %d. Chunk: %d\n", tid, chunk);
36
37 begin = tid*chunk;
38 end = begin + chunk;
39
40 for (int i=begin; i<end; i++)
41 array[i] = tid;
42
43 return 0;
44 }
```

# Parallel region

## #omp parallel firstprivate(list of variables)

- In general, private variables are not initialised at the beginning of the parallel region.
- However, the variables in the list of `firstprivate` are initialised.

```
1 # include <omp.h>
2 # include <stdio.h>
3
4 int main () { //Warning! it might show race conditions
5 omp_set_num_threads(4);
6
7 int a = 2;
8 int b = 3;
9
10 # pragma omp parallel firstprivate(a) shared(b)
11 {
12 int id = omp_get_thread_num();
13
14 printf("Thread %d. a: %d. b: %d\n",id ,a,b);
15 a++; b++;
16 printf("Thread %d. a++: %d. b++: %d\n",id ,a,b);
17 }
18
19 printf("Master. a: %d. b: %d\n",a,b);
20
21 return 0;
22 }
```

# Parallel region

## #omp parallel reduction(operator : list of variables)

- Produces a single value from associative operations (+, \*, max, min, and, or).

```
1 # include <omp.h>
2 # include <stdio.h>
3 # define R 2
4 # define C 3
5
6 int main () {
7 int b = 0; //Try with b = 100 instead!
8 int a[R][C] = {{1,2,3},{4,5,6}};
9
10 omp_set_num_threads(2);
11
12 # pragma omp parallel reduction(+:b)
13 {
14 int id = omp_get_thread_num();
15 //Try here setting b = 100
16
17 for (int i=0;i<C;i++){
18 b = b + a[id][i];
19 }
20 printf("Thread %d working on line %d of a. b = %d\n",id,id,b);
21 }
22
23 printf("Back to master. b = %d\n",b);
24 return 0;
25 }
```