

Shared-memory programming : OpenMP (II)

ING2-GSI-MI Architecture et Programmation Parallèle

Juan Angel Lorenzo del Castillo
juan-angel.lorenzo-del-castillo@cyu.fr

CY Cergy Paris Université
2023-2024



Table of Contents

1 OpenMP (cont.)

Table of Contents

1 OpenMP (cont.)

OpenMP directives

Most relevant OpenMP directives

- Parallel regions construction
 - ▶ `parallel`
- Work sharing
 - ▶ `for`, `sections`, `single`.
- Synchronisation
 - ▶ `master`, `critical`, `atomic`, `barrier`.
- Task management
 - ▶ `task`, `taskwait`.

There are more...

OpenMP directives

Most relevant OpenMP directives

- Parallel regions construction
 - ▶ `parallel`
- Work sharing
 - ▶ **`for, sections, single.`**
- Synchronisation
 - ▶ `master, critical, atomic, barrier.`
- Task management
 - ▶ `task, taskwait.`

There are more...

Today's class

Work sharing

#pragma omp for

Syntax:

```
#pragma omp for [clauses]  
for loop
```

- A parallel *for* loop divides up the iterations of the loop between threads.
- With no additional clauses, the *for* directive will usually partition the iterations as equally as possible between the threads.
- The iterations of the loop will be executed in parallel by the thread team.
- The parallel loop index variable is private.
- No new threads are created: the iterations of the loop are divided up between the existing threads from the thread team.
- **Implicit barrier** at the end of the loop unless the **nowait** clause was used.

Work sharing

#pragma omp for

Syntaxe :

```
#pragma omp for [clauses]  
for loop
```

Allowed clauses:

- private(list of variables)
- firstprivate(list of variables)
- lastprivate(list of variables)
- reduction(operator: list of variables)
- schedule(type[,block size])
- ordered
- collapse(positive value)
- nowait

Work sharing

#pragma omp for

```

1  # include <omp.h>
2  # include <stdio.h>
3
4  int main () {
5
6      int i,id;
7      const int SIZE = 12;
8      int A[SIZE];
9
10     omp_set_num_threads(4);
11
12     #pragma omp parallel
13     {
14         #pragma omp for private(id)
15         for (i=0; i<SIZE;i++)
16         {
17             id=omp_get_thread_num();
18             A[i] = id;
19         }
20     }
21
22     for (i=0; i<SIZE; i++) {
23         printf(" | %d",A[i]);
24     }
25     printf(" | \n");
26
27     return 0;
28 }
```

```

1  # include <omp.h>
2  # include <stdio.h>
3
4  int main () {
5
6      int i,id;
7      const int SIZE = 12;
8      int A[SIZE];
9
10     omp_set_num_threads(4);
11
12     #pragma omp parallel for private(id)
13     for (i=0; i<SIZE;i++)
14     {
15         id=omp_get_thread_num();
16         A[i] = id;
17     }
18
19
20     for (i=0; i<SIZE; i++) {
21         printf(" | %d",A[i]);
22     }
23     printf(" | \n");
24
25     return 0;
26 }
```


Work sharing

#pragma omp for

What is the difference between these two programs?

```

1 # include <omp.h>
2 # include <stdio.h>
3
4 int main ()
5 {
6     int tid;
7     omp_set_num_threads(4);
8
9     #pragma omp parallel private(tid)
10 {
11     for (int i=0;i<10;i++)
12     {
13         tid = omp_get_thread_num();
14         printf("Thread %d i = %d\n",tid,i);
15     }
16 }
17 return 0;
18 }
```

```

1 # include <omp.h>
2 # include <stdio.h>
3
4 int main ()
5 {
6     int tid;
7     omp_set_num_threads(4);
8
9     #pragma omp parallel for private(tid)
10 for (int i=0;i<10;i++)
11 {
12     tid = omp_get_thread_num();
13     printf("Thread %d i = %d\n",tid,i);
14 }
15
16 return 0;
17 }
```

Work sharing

#pragma omp for private(list of variables)

#pragma omp for firstprivate(list of variables)

#pragma omp for reduction(operator: list of variables)

- Already seen in the *Parallel region* section.

Work sharing

#pragma omp for lastprivate(list of variables)

- The variables of the list are declared **private**.
 - ▶ Reminder: private variables are undefined before and after the parallel *for* loop.
- The variables of the list are updated with their last value at the end of the parallel region.
- This last value is the value given by the thread that has executed the last iteration of the loop.

Work sharing

#pragma omp for lastprivate(list of variables)

Example:

```

1  #include <omp.h>
2  #include <stdio.h>
3
4  int show(int , int);
5
6  int main () {
7      int a = 2;
8      int tid;
9
10     omp_set_num_threads(4);
11
12     # pragma omp parallel for
13         firstprivate(a) lastprivate(a)
14         private(tid)
15     for (int i = 0; i < 5; i++)
16     {
17         tid = omp_get_thread_num();
18         a++;
19         show(tid , a);
20     }
21
22     printf("Master. a: %d\n",a);
23     return 0;
24 }

```

```

24 int show(int id , int a)
25 {
26     // ignore this pragma for now
27     #pragma omp critical
28     {
29         printf("Thread %d. a: %d\n",id ,a);
30     }
31     return 0;
32 }

```

Work sharing

#pragma omp for schedule(type[,block size])

- Gives a variety of options for specifying which loops iterations are executed by which thread.
- Important for load balancing purposes (i.e. performance).
- Types:
 - ▶ schedule(static[,chunk])
 - ▶ schedule(dynamic[,chunk])
 - ▶ schedule(guided[,chunk])
 - ▶ schedule(runtime)
 - ▶ schedule(auto)

Work sharing

#pragma omp for schedule(static[,chunk])

- With no `chunk` specified, the iteration space is divided into (approximately) equal chunks ($\text{num_it}/\text{num_threads}$), and one `chunk` is assigned to each thread (**block schedule**).
- If `chunk` is specified, the iteration space is divided into chunks, each of `chunk` iterations, and the chunks are assigned cyclically to each thread (**block cyclic schedule**)

#pragma omp for schedule(dynamic[,chunk])

- It divides the iteration space up into chunks of size `chunk` (like *static*), but chunks are assigned to threads on a *first-come-first-served* basis.
- i.e. as a thread finishes a chunk, it is assigned the next chunk in the list.
- When no `chunk` is specified, it defaults to 1.

Work sharing

`#pragma omp for schedule(guided[,chunk])`

- similar to *dynamic*, but the chunks start off large ($\text{num_it}/\text{num_threads}$) and get smaller exponentially.
- The size of the next chunk is (roughly) the number of remaining iterations divided by the number of threads.
- The `chunk` specifies the minimum size of the chunks.
- When no `chunk` is specified, it defaults to 1.

`#pragma omp for schedule(runtime)`

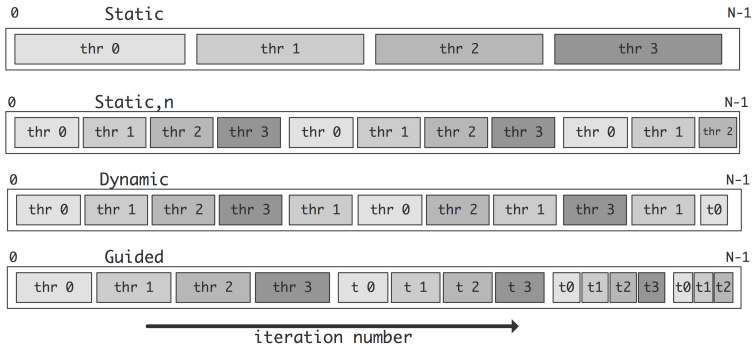
- It defers the choice of schedule to run time, when it is determined by the value of the environment variable `OMP_SCHEDULE`:

▶ `export OMP_SCHEDULE = "guided,4"`

`#pragma omp for schedule(auto)`

- The choice of scheduling is left to the compiler.

Work sharing



Source : utexas.edu

Work sharing

When to choose which schedule:

- `Static` is the best for already-load balanced loops (least overhead).
- `Static, n` is good for loops with mild or smooth load imbalance, but can induce *false sharing* (to be seen in the exercises).
- `Dynamic` is useful if iterations have widely varying loads, but ruins data locality.
- `Guided` is often less expensive than `Dynamic`, but beware of loops where the first iterations are the most expensive.
- `Runtime` can be used for convenient experimentation of the different scheduling techniques.

Work sharing

#pragma omp for collapse(positive value)

When we have perfectly nested loops we can collapse inner loops. Collapsing a loop:

- The loop needs to be perfectly nested.
- The loop needs to have rectangular iteration space.
- Makes iteration space larger: iterations from all loops are grouped to make a single iteration space that will be shared among the threads.
- Less synchronisation needed than nested parallel loops.

```
#pragma omp for collapse(2)
for(i=0; i<10; i++)
  for(j=0; j<10; j++)
    myfunction(i, j);
```

\Leftrightarrow
(similar to)

```
#pragma omp for
for(i=0; i<100; i++)
  myfunction(i/10, i%10);
```

Work sharing

#pragma omp for nowait

- Synchronisation directive to suppress the implicit barriers at the end of `for`, `sections` and `single` directives (barriers are expensive).

```

1  # include <omp.h>
2  # include <stdio.h>
3
4  int main () {
5
6      int i;
7      const int SIZE = 12;
8      int A[SIZE];
9      int B[SIZE];
10     int C[SIZE];
11     int D[SIZE];
12
13     omp_set_num_threads(4);
14
15     // Initialisation
16     for (i = 0; i < SIZE; i++)
17     {
18         A[i] = 100;
19         B[i] = 100;
20         C[i] = 100;
21         D[i] = 100;
22     }
23
24     #pragma omp parallel private(i) shared(A,B,C,D)
25     {
26         #pragma omp for nowait
27         for (i = 0; i < SIZE-1; i++)
28         {
29             B[i] = (A[i] + A[i+1])/2;
30         }
31
32         #pragma omp for
33         for (i = 0; i < SIZE; i++)
34         {
35             D[i] = 1/C[i];
36         }
37     }
38
39     return 0;
40 }

```

Work sharing

#pragma omp sections

Syntax:

```
#pragma omp sections [clauses]
{
    #pragma omp section
        structured block
    #pragma omp section
        structured block
    ...
}
```

- Allows separate blocks of code to be executed in parallel (e.g. several independent subroutines)
- **Implicit barrier** at the end, unless `nowait` is specified.
- Rarely used, except with nested parallelism.

Work sharing

#pragma omp sections

Syntax (shorthand form):

```
#pragma omp [parallel] sections [clauses]
{
    #pragma omp section
        structured block
    #pragma omp section
        structured block
    ...
}
```

Allowed clauses:

- private(liste de variables)
- firstprivate(liste de variables)
- lastprivate(liste de variables)
- reduction(operateur : liste de variables)
- nowait

Work sharing

#pragma omp single

Syntax:

```
#pragma omp single [clauses]
{
    structured block
}
```

- The block of code will be executed by a single thread only.
- The first thread to reach the `single` directive will execute the block.
- Other threads wait until the block has been executed.
- **Implicit barrier** at the end unless `nowait` is specified.

Work sharing

#pragma omp single

Syntax:

```
#pragma omp single [clauses]
{
    structured block
}
```

Allowed clauses:

- `private(list of variables)`
- `firstprivate(list of variables)`
- `copyprivate(list of variables)`
 - ▶ At the end of the `single` directive, the value of the private variable(s) is copied to the private variable(s) of the other threads.
 - ▶ Incompatible with `nowait`.
- `nowait`