

Performance optimisation in parallel systems

Juan Ángel Lorenzo del Castillo

CY Cergy Paris Université

ING2-GSI-MI Architecture et Programmation Parallèle

2023 - 2024



juan-angel.lorenzo-del-castillo@cyu.fr

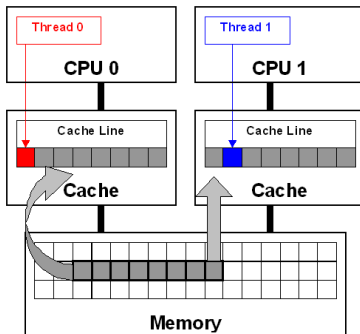
Table of Contents

- 1 Performance optimisation in shared-memory systems
- 2 Performance optimisation in distributed-memory systems

Table of Contents

- 1 Performance optimisation in shared-memory systems
- 2 Performance optimisation in distributed-memory systems

False sharing



Source : Intel

- It occurs when threads on different processors modify variables that reside on the same cache line. This invalidates the cache line and forces an update, which hurts performance.

False sharing

Example

```
double sum=0.0, sum_local[NUM_THREADS];

#pragma omp parallel num_threads(NUM_THREADS)
{
    int tid = omp_get_thread_num();
    sum_local[tid] = 0.0;

    #pragma omp for
    for (i = 0; i < N; i++)
        sum_local[tid] += x[i] * y[i];

    #pragma omp atomic
    sum += sum_local[tid];
}
```

Techniques to reduce False sharing

- Use private variables whenever it is possible.
- Modify data or iterations-to-thread affinity.
- Give more work per *chunk* to each thread.
- Use a more intelligent compiler.
- Its impact will be less noticeable on larger problems.

Table of Contents

- 1 Performance optimisation in shared-memory systems
- 2 Performance optimisation in distributed-memory systems

Improving locality in distributed-memory systems

■ Cannon's Matrix Multiplication Algorithm

- ▷ Assume a matrix of size p that is a perfect square
- ▷ Each processor gets a $n/\sqrt{p} \times n/\sqrt{p}$ chunk of data
- ▷ Organise processors into rows and columns, keeping data locality for each processor

$p(0,0)$	$p(0,1)$	$p(0,2)$
$p(1,0)$	$p(1,1)$	$p(1,2)$
$p(2,0)$	$p(2,1)$	$p(2,2)$

Source : Scott B. Baden / Jim Demmel

Improving locality in distributed-memory systems

■ Cannon's Matrix Multiplication Algorithm

- ▷ Move data incrementally in \sqrt{p} phases
- ▷ Circulate each chunk of data among processors within a row or column
- ▷ Consider iteration $i = 1, j = 2$:

$$C[1,2] = A[1,0]*B[0,2] + A[1,1]*B[1,2] + A[1,2]*B[2,2]$$

A(0,0)	A(0,1)	A(0,2)
A(1,0)	A(1,1)	A(1,2)
A(2,0)	A(2,1)	A(2,2)

B(0,0)	B(0,1)	B(0,2)
B(1,0)	B(1,1)	B(1,2)
B(2,0)	B(2,1)	B(2,2)

Scott B. Baden / Jim Demmel

Improving locality in distributed-memory systems

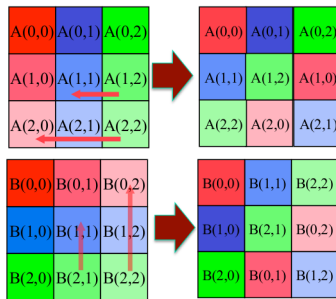
■ Cannon's Matrix Multiplication Algorithm

$$C[1,2] = \mathbf{A}[1,0] * \mathbf{B}[0,2] + A[1,1] * B[1,2] + A[1,2] * B[2,2]$$

- Initially we want $A[1,0]$ and $B[0,2]$ to reside on the same processor as $C[1,2]$.

So we first skew the matrices for everything to line up:

- Shift each row i by i columns to the left.
- Communication wraps around.
- Same for each column.



Source : Scott B. Baden / Jim Demmel

Improving locality in distributed-memory systems

■ Cannon's Matrix Multiplication Algorithm

$$C[1,2] = A[1,0]*B[0,2] + \mathbf{A[1,1]*B[1,2]} + A[1,2]*B[2,2]$$

2. Then we shift rows and columns so the next pair of values $A[1,1]$ and $B[1,2]$ line up.

We circularly shift:

- each row by 1 column to the left.
- each column by 1 row to the "north".

Each processor calculates the product of the two local sub-matrices adding into the accumulated sum.

A(0,1)	A(0,2)	A(0,0)
A(1,2)	A(1,0)	A(1,1)
A(2,0)	A(2,1)	A(2,2)



B(1,0)	B(2,1)	B(0,2)
B(2,0)	B(0,1)	B(1,2)
B(0,0)	B(1,1)	B(2,2)



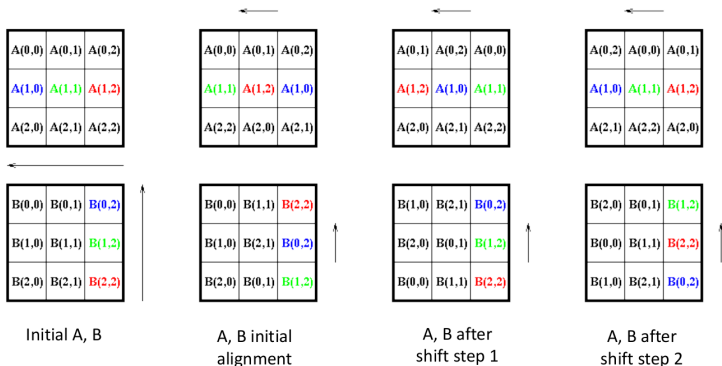
Source : Scott B. Baden / Jim Demmel

Improving locality in distributed-memory systems

■ Cannon's Matrix Multiplication Algorithm

$$C[1,2] = A[1,0]*B[0,2] + A[1,1]*B[1,2] + A[1,2]*B[2,2]$$

■ Summary:



Source : Scott B. Baden / Jim Demmel

Improving locality in distributed-memory systems

■ Limitations of Cannon's Algorithm:

- ▷ p must be a perfect square
- ▷ A and B must be square, and evenly divisible by \sqrt{p}